

Triangle Rasterization

Computer Graphics
COMP 770 (236)
Spring 2007

Instructor: Brandon Lloyd

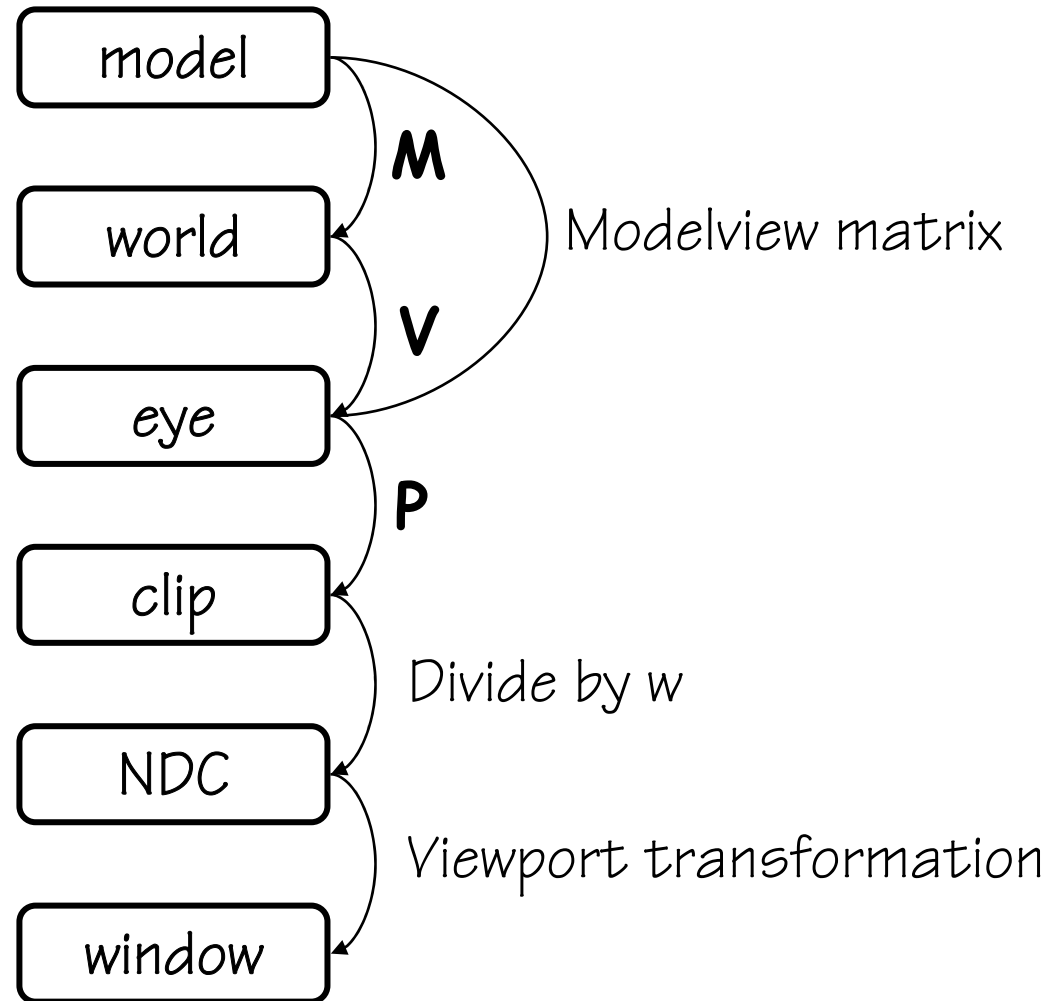
From last time...

- Lines and planes
- Culling
 - View frustum culling
 - Back-face culling
 - Occlusion culling
 - Hierarchical culling
- Clipping

Topics for today


- Quick review of coordinate systems
- Motivation
 - What is rasterization?
 - Why triangles?
- Rasterization
 - Scan-line
 - Edge equations
- Interpolation
- Beyond triangles

Coordinate systems



Primitive rasterization


- Rasterization converts vertex representation to pixel representation
 - *Coverage determination* – Computes which pixels (samples) belong to an “ideal” analytical primitive
 - *Parameter interpolation* – Computes parameters at covered pixels from parameters associated with primitive vertices
- Coverage is a 2-D sampling problem
- Possible coverage criteria:
 - distance of the primitive to sample point (often used with lines)
 - Percent coverage of a pixel (used to be popular)
 - Sample is inside the primitive (assuming it is closed)



Why triangles?

■ Triangles are simple

- minimal representation for a surface element (3 points or 3 edge equations)
- triangles are linear (makes computations easier)





$$\mathcal{T} = (\dot{v}_0, \dot{v}_1, \dot{v}_2)$$

$$\mathcal{T} = (\bar{e}_0, \bar{e}_1, \bar{e}_2)$$

Why triangles?

- Triangles are **convex**
- What does it mean to be a convex?





An object is convex if and only if any line segment connecting two points on its boundary is contained entirely within the object or one of its boundaries.

- Why is convexity important?

Regardless of a triangle's orientation on the screen a given scan line will contain only a single segment or *span* of that triangle.

Why triangles?


- Arbitrary polygons can be decomposed into triangles



- Decomposing a convex n -sided polygon is trivial
 - Suppose the polygon has ordered vertices $\{v_0, v_1, \dots, v_n\}$
 - It can be decomposed into triangles $\{(v_0, v_1, v_2), (v_0, v_2, v_3), (v_0, v_i, v_{i+1}), \dots, (v_0, v_{n-1}, v_n)\}$.
- Decomposing a non-convex polygon is non-trivial
 - sometimes have to introduce new vertices


Why triangles?

- Triangles can approximate any 2-dimensional shape (or 3D surface)
- Polygons are a locally linear (planar) approximation.
- Improve the quality of fit by increasing the number edges or faces.




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges




Scanline triangle rasterizer

- Walk along edges one scanline at a time
- Rasterize spans between edges





Fractional offsets



- Straightforward to interpolate values (e.g. colors) along the edges, but must be careful when offsetting from the edge to the pixel's center.

Scanline rasterizing entire scenes

- Sort all edges by start scanline into the Inactive Edge Table (IET)
- Move edges intersected by current scanline from IET to Active Edge Table (AET)
- Compute spans between active edges
- Sort spans by starting x
- Rasterize visible span segments
- Remove edges from AET when they no longer intersect the current scanline



Scanline rasterization

■ Advantages:


- Can be made quite fast
- Low memory usage for smallish scenes
- Don't need full 2D z-buffer (can use 1D z-buffer on the scanline)

■ Disadvantages:

- Doesn't scale well to large scenes
- Have to worry about fractional offsets
- Lots of special cases


Rasterizing with edge equations

- Compute edge equations from vertices
- Compute interpolation equations from vertex parameters
- Traverse pixels evaluating the edge equations
- Draw pixels for which all edge equations are positive
- Interpolate parameters at pixels



Edge equation coefficients

- The cross product between 2 homogeneous points generates the line between them



$$\begin{aligned}\bar{e} &= \dot{v}_0 \times \dot{v}_1 \\ &= [x_0 \quad y_0 \quad 1]^t \times [x_1 \quad y_1 \quad 1]^t \\ &= [(y_0 - y_1) \quad (x_1 - x_0) \quad (x_0 y_1 - x_1 y_0)] \\ &\quad \quad \quad A \quad \quad \quad B \quad \quad \quad C\end{aligned}$$

$$E(x, y) = Ax + By + C$$

- A pixel at (x, y) is “inside” an edge if $E(x, y) > 0$

Numerical precision

- Subtraction of two nearly equal floating point numbers results in catastrophic cancellation which leaves only a few significant bits

$$1.234 \times 10^3 - 1.233 \times 10^3 = 1.000 \times 10^0$$

- When $x_0 y_1 \approx x_1 y_0$ computing $C = x_0 y_1 - x_1 y_0$ can result in loss of precision
- Reformulate C coefficient:

$$C = -\frac{A(x_0 + x_1) + B(y_0 + y_1)}{2}$$


Triangle area

$$\begin{aligned}\text{Area} &= \frac{1}{2} \det \begin{bmatrix} x_0 & x_1 & x_2 \\ y_0 & y_1 & y_2 \\ 1 & 1 & 1 \end{bmatrix} \\ &= \frac{1}{2} ((x_1 y_2 - x_2 y_1) - (x_0 y_2 - x_2 y_0) + (x_0 y_1 - x_1 y_0)) \\ &= \frac{1}{2} (C_0 + C_1 + C_2)\end{aligned}$$

- $\text{Area} = 0$ means that the triangle is not visible
- $\text{Area} < 0$ means the triangle is back facing:
 - Reject triangle if performing back-face culling
 - Otherwise, flip edge equations by multiplying by -1

Shared edges

- Suppose two triangles share an edge. Which covers the pixel when the edge passes through the sample ($E(x,y)=0$)?
- Both
 - pixel color becomes dependent on order of triangle rendering
 - creates problems when rendering transparent objects - “double hitting”
- Neither
 - Missing pixels create holes in otherwise solid surface
- We need a consistent tie-breaker!




Shared edges


- A common tie-breaker:

$$\text{bool } t = \begin{cases} A > 0 & \text{if } A \neq 0 \\ B > 0 & \text{otherwise} \end{cases}$$


- Coverage determination becomes
if($E(x,y) > 0 \parallel (E(x,y) == 0 \ \&\& \ t)$)
pixel is covered



Shared vertices



- Use “inclusion direction” as a tie breaker.
- Any direction can be used



- Snap vertices to subpixel grid and displace so that no vertex can be at the pixel center

Other benefits of snapping to subpixel grid

■ Simplicity

- can use fixed-point arithmetic can be used (integer operations)

■ Robustness


- With sufficient bits, edge equations and areas can be computed exactly

■ Quality

- Smoother animation than if we snapped to the pixel grid

Interpolating parameters

- Specify a parameter, say redness (r) at each vertex of the triangle.
- Linear interpolation creates a planar function



$$r(x,y) = Ax + By + C$$

Solving for interpolation equation

- Given the redness of the three vertices, we can set up the following linear system:

$$\begin{bmatrix} r_0 & r_1 & r_2 \end{bmatrix} = \begin{bmatrix} A_r & B_r & C_r \end{bmatrix} \begin{bmatrix} x_0 & x_1 & x_2 \\ y_0 & y_1 & y_2 \\ 1 & 1 & 1 \end{bmatrix}$$

with the solution:

$$\begin{bmatrix} A_r & B_r & C_r \end{bmatrix} = \begin{bmatrix} r_0 & r_1 & r_2 \end{bmatrix} \frac{\begin{bmatrix} (y_1 - y_2) & (x_2 - x_1) & (x_1 y_2 - x_2 y_1) \\ (y_0 - y_2) & (x_2 - x_0) & (x_0 y_2 - x_2 y_0) \\ (y_0 - y_1) & (x_1 - x_0) & (x_0 y_1 - x_1 y_0) \end{bmatrix}}{\det \begin{bmatrix} x_0 & x_1 & x_2 \\ y_0 & y_1 & y_2 \\ 1 & 1 & 1 \end{bmatrix}}$$

Interpolation equation

- The parameter plane equation is just a linear combination of the edge equations


$$\begin{bmatrix} A_r & B_r & C_r \end{bmatrix} = \frac{1}{2 \cdot \text{area}} \begin{bmatrix} r_0 & r_1 & r_2 \end{bmatrix} \begin{bmatrix} \bar{e}_0 \\ \bar{e}_1 \\ \bar{e}_2 \end{bmatrix}$$

Extra work to interpolate a parameter:

- Transform parameter vector
- Compute one interpolation equation per pixel per parameter

Z-Buffering

- When rendering multiple triangles we need to determine which triangles are visible
- Use z-buffer to resolve visibility
 - stores the depth at each pixel
- Initialize z-buffer to 1
 - Post-perspective z values lie between 0 and 1
- Linearly interpolate depth (z_{tri}) across triangles
 - Why can we do this?
- If $z_{\text{tri}}(x,y) < \text{zBuffer}[x][y]$
write to pixel at (x,y)
 $\text{zBuffer}[x][y] = z_{\text{tri}}(x,y)$





A simple three dimensional scene



Z-buffer representation

Traversing pixels

- Free to traverse pixels how we please
 - Edge and interpolation equations can be computed at any point
- Try to minimize work
 - Restrict traversal to primitive bounding box
 - Zig-zag traversal avoids empty pixels
 - Hierarchical traversal
 - Knock out tiles of pixels (say 4x4) at a time
 - Test corners of tiles against equations
 - Test individual pixels of tiles not entirely inside or outside



Incremental update of linear equations

- Some computation can be saved by updating the edge and interpolation equations incrementally:

$$E(x, y) = Ax + By + C$$

$$\begin{aligned} E(x + \Delta, y) &= A(x + \Delta) + By + C \\ &= E(x, y) + A \cdot \Delta \end{aligned}$$

$$\begin{aligned} E(x, y + \Delta) &= Ax + B(y + \Delta) + C \\ &= E(x, y) + B \cdot \Delta \end{aligned}$$

- Equations can be updated with a single addition!

Triangle setup

- Compute edge equations
 - 3 cross products
- Compute triangle area
 - A few additions
- Cull zero area and back-facing triangles and/or flip edge equations
- Compute interpolation equations
 - Matrix/vector product per parameter

A Post-Triangle World?

- Are triangles really the best rendering primitive?


100,000,000 primitives

2,000,000 pixels

5 primitives/pixel

(Assuming that primitives are uniformly distributed over screen and only 10% are visible.)

- Cost to render a single triangle
 - specify 3 vertices
 - compute 3 edge equations
 - evaluate equations one



Models of this magnitude are being built today. The leading and most ambitious work in this area is Stanford's "Digital Michelangelo Project".

Point-cloud rendering

- Points have been proposed as a rendering primitive to address this problem.
- Key Attributes:
 - Hierarchy
 - Incremental refinement
 - Compact representation (differential encoding)



130,712 Splats, 132 mS



259,975 Splats, 215 mS



1,017,149 Splats, 722 mS



14,835,967 Splats, 8308 mS

Next time

- Texture mapping
- Barycentric coordinates
- Perspective-correct interpolation
- Texture filtering